


THE MODERN AI SKILLS HANDBOOK

A PRACTICAL GUIDE TO AI, PROMPTING, LLMS, AGENTS,
AUTOMATION, VIDEO CREATION, VIBE CODING, +
PYTHON, EXPLORATORY DATA ANALYSIS
MACHINE LEARNING



CONTENTS

- 01 AI FOUNDATIONS**
- 02 PROMPT ENGINEERING**
- 03 LARGE LANGUAGE MODELS**
- 04 AI AGENTS**
- 05 AGENTIC AI**
- 06 AUTOMATION**
- 07 AI VIDEO MAKING**
- 08 VIBE CODING**
- 09 PYTHON**
- 10 EXPLORATORY DATA ANALYSIS**
- 11 MACHINE LEARNING**
- 12 FUTURE OF AI**
- 13 ACTIVITIES**
- 14 SUMMARY**



INTRODUCTION

Most people encounter AI as a collection of disconnected tools a chatbot here, an image generator there, an automation script somewhere 'else. What's missing is a coherent picture of how these pieces relate to one another and why understanding that relationship matters.

This handbook takes a different approach. It treats the AI landscape as a connected set of skills, each one building on the previous, each one enabling the next. By the end, you won't just know what these technologies are you'll understand how they work together and where you fit into that picture.

HOW THIS GUIDE IS STRUCTURED

The topics in this handbook follow a deliberate sequence. We begin with AI foundations the underlying concepts that make everything else interpretable. From there we move into prompt engineering, because your ability to communicate with AI shapes everything you can do with it.


Once you understand how to talk to AI, we examine what you're actually talking to: large language models, their strengths, and their genuine limitations. That leads naturally into agents and agentic systems AI that doesn't just answer questions but takes action and plans across multiple steps.

The second half of the guide shifts toward practical application: automation workflows, AI-assisted video production, the emerging practice of vibe coding, Python, and finally two foundational data science disciplines - Exploratory Data Analysis and Machine Learning that complete the modern AI practitioner's toolkit.

WHO THIS GUIDE IS FOR

Anyone who wants to understand modern AI clearly without jargon, without hype, and without needing a computer science degree. Whether you're a professional a creator, a student, or simply someone who wants to engage intelligently with a fast-moving field, this is Your starting point.

One note on scope: This guide introduces concepts and builds understanding. It is not a technical manual or a complete course. Think of it as the orientation before the deeper work begins the map before the territory.



01 AI FOUNDATIONS

Artificial intelligence is at its core, software that performs tasks which typically require some form of judgment or pattern recognition. That definition is deliberately simple because the popular understanding of AI robots, sci-fi, superintelligence ends to obscure what the technology actually does.

In practical terms, AI systems are trained on large amounts of data, develop internal representations of patterns within that data and use those patterns to make predictions or generate outputs. A spam filter, a recommendation algorithm, and a language model are all AI systems, even though they look and behave very differently.

TYPES OF AI SYSTEMS

NOT ALL AI IS THE SAME. THE FIELD CONTAINS SEVERAL DISTINCT FAMILIES OF SYSTEMS, EACH SUITED TO DIFFERENT PROBLEMS.

Type	Description
Rule-based AI	Follows explicit logic trees. Reliable and rigid. It was commonly used in early chatbots and decision systems.
Machine Learning	Learns patterns from data without being explicitly programmed. It is the dominant approach in use today.
Deep Learning	A subset of machine learning that uses neural networks with many layers. It powers image recognition, speech processing, and language models.
Generative AI	Generates new content — such as text, images, code, and audio — rather than just classifying existing data.
Reinforcement Learning	Learns through trial and error by optimizing for rewards. It is used in game-playing AI and robotics.




TRAINING

AI models are trained exposed to large datasets and adjusted over many iterations until they become reliably useful. A model behavior is a direct consequence of what it was trained on and how that training was designed.




INFERENCE

Once trained, a model makes predictions or generates outputs from new inputs. This is called inference. When you type a message to a chatbot, the model runs inference to produce a response.




PARAMETERS

Inside a neural network, parameters are numerical values that determine how inputs get processed. Larger models have more parameters and can generally handle more complex tasks but not always. Architecture and training quality matter as much as raw scale.



CONTEXT WINDOW

Most modern language models work with a context window the amount of text they can see at once when generating a response. Think of it as working memory. Longer context windows allow more nuanced, coherent outputs over extended conversations or documents.



COMMON MISCONCEPTION

AI does not think. It processes patterns in data to produce statistically likely outputs. This is sophisticated and often useful but it is not cognition in any human sense. Understanding this distinction helps you use AI more effectively and set realistic expectations.

02 PROMPT ENGINEERING

COMMUNICATING WITH AI EFFECTIVELY

If AI is a capable collaborator, prompts are how you direct that collaboration. Prompt engineering is the practice of designing inputs to language models in ways that produce the most useful, accurate, and relevant outputs. It is closer to clear thinking and clear writing than to programming.

Language models do not read your mind. Two prompts asking for the same thing in different ways can produce dramatically different results. The model did not change the input did. This means improving your results does not require understanding a model's internals; it requires communicating clearly and deliberately.

ANATOMY OF A GOOD PROMPT

Element	Description
Context	Background the model needs: who you are, what you are working on, why it matters.
Task	What you want produced. Be specific about format, length, tone, and audience.
Constraints	What to avoid or include. Restrictions help the model narrow its output.
Examples	Showing the model what 'good' looks like dramatically improves results.

WEAK VS STRONG PROMPT

Weak Prompt	Strong Prompt
Tell me about marketing.	I'm a freelance graphic designer launching a one-person studio. Write three short LinkedIn post ideas (under 150 words each) that highlight my services without sounding promotional. Tone: professional but approachable.


CORE TECHNIQUES

Role Assignment: Asking the model to adopt a role shapes its perspective. 'You are an experienced project manager reviewing this proposal' produces different feedback than the same question without that framing.

Step-by-Step Reasoning: For complex problems, ask the model to reason through them before reaching a conclusion. This produces more reliable, traceable answers.

Iteration: Treat prompting as a conversation. Review the output, identify what is off, and refine. Most professional prompt workflows involve multiple rounds of adjustment.

Format Specification: If you want a table, ask for a table. If you want 200 words, say so. Models adapt to format instructions reliably when those instructions are clear.



KEY INSIGHT

Prompt engineering is not about manipulating AI. It is about communicating precisely. Every skill that makes you a clearer thinker and writer makes you a better prompt engineer.

03 LARGE LANGUAGE MODELS

THE ENGINES BEHIND MODERN AI



Large Language Models LLMs are the technology behind most AI tools people interact with daily: chatbots, writing assistants, code helpers, document analysers. At the most basic level, an LLM predicts what text should come next given what it has already seen. It was trained on vast amounts of written material books, articles, code, conversations and developed an internal representation of how language works, how ideas relate, and how topics are discussed.

When you send a message, the model does not look up an answer in a database. It generates a response token by token, each one chosen based on probability distributions learned during training. This is why LLMs can write poetry, explain concepts, debug code, and translate languages they have absorbed patterns across all those domains.

REAL-WORLD APPLICATIONS

Application	Description
Writing assistance	Drafting, editing, summarising, and improving written content across formats.
Code generation	Writing, explaining, and debugging code across multiple programming languages.
Document analysis	Extracting key information, answering questions, and summarising long texts.
Customer service	Powering conversational interfaces that handle common queries at scale.
Research synthesis	Organising and summarising information from multiple sources.
Language translation	High-quality translation across dozens of language pairs.

STRENGTHS AND LIMITATIONS


LLMs are genuinely good at tasks involving language: explaining complex ideas, generating structured content, adapting tone and style, and synthesising information across a broad knowledge base.

They also have real limitations: a knowledge cutoff means they are unaware of events after their training ended; they can produce confident-sounding but incorrect information (hallucination); they struggle with precise arithmetic; and they have no memory between conversations unless that memory is explicitly provided.


KEY INSIGHT

The right approach to LLM limitations is not avoidance it is design. Structure your workflows to verify factual claims, provide relevant context in your prompts, and treat LLM output as a skilled first draft rather than a final product.

HOW LLMS WORK (SIMPLIFIED)



WHAT LLMS CAN DO



LIMITATIONS



04 AI AGENTS

AI THAT TAKES ACTION



A standard language model takes an input, produces an output, and stops. An AI agent does something more: it uses language model reasoning to take action, evaluate results, and continue until a goal is reached. This distinction from generating to doing is the defining feature of agent-based AI.

When you ask a chatbot to write a summary, it reads your input and produces an output. When an agent is given the same task, it might search for the article, retrieve relevant sections, evaluate their importance, draft a summary, check the word count, and revise all without being prompted at each step.

CORE COMPONENTS OF AN AI AGENT

Component	Description
Perception	The agent takes in information — text, data, tool outputs — to understand its current situation.
Reasoning	Using an LLM, the agent interprets the situation and decides what to do next.
Action	The agent executes a step: searching the web, writing code, calling an API, updating a file.
Memory	Short-term context from the current task, and sometimes longer-term storage for persistence.
Evaluation	The agent checks whether its output meets the goal and adjusts if needed.

REAL-WORLD AGENT EXAMPLES

- ◆ Research agents that gather information from multiple sources and synthesise it into structured reports.
- ◆ Coding agents that write code, run it, observe errors, and iteratively debug until the programme works.
- ◆ Customer support agents that understand a complaint, look up account history, and draft a resolution.
- ◆ Data agents that pull from a database, run analysis, generate visualisations and summarise findings.

KEY INSIGHT

Agents are only as good as the tools and permissions they are given.
Building effective agent workflows means thinking carefully about what the agent needs to access and what it should not.



05 AGENTIC AI


AUTONOMOUS, MULTI-STEP SYSTEMS

Agentic AI describes systems in which AI agents operate with greater autonomy, collaborate with each other and manage complex multi-step goals with minimal human intervention at each stage. Three properties distinguish agentic systems from simpler AI tools.

- ◆ **AUTONOMY:** The system pursues goals across multiple steps without requiring human approval at each stage.
- ◆ **PERSISTENCE:** The system maintains context and state across an extended workflow, not just a single exchange.
- ◆ **ADAPTATION:** When something unexpected happens, the system adjusts its approach rather than stopping.

MULTI-STEP PLANNING

Agentic AI systems typically use a plan-then-execute model. Before taking any action, the system breaks the goal into sub-tasks, orders them logically identifies dependencies, and allocates effort. Execution follows the plan with the system checking progress and revising the plan if needed.



MULTI-AGENT COLLABORATION



Some of the most capable agentic systems involve not one agent but several, each with a defined role. A lead agent decomposes a goal and delegates sub tasks to specialist agents one for research, one for writing, one for factchecking. The lead agent then integrates their outputs into a final result Software development, scientific research, business operations, and creative production are all fields where agentic systems are being piloted not as replacements for human judgment, but as multipliers of human capacity.

FUTURE POTENTIAL

The most significant applications of agentic AI are still emerging. Software development, scientific research, business operations, and creative production are all fields where agentic systems are being piloted not as replacements for human judgement, but as multipliers of human capacity.

EXAMPLE WORKFLOW

Goal: Prepare a competitive analysis report. An agentic system might plan: identify competitors research each one extract key metrics compare against a framework draft the report format for presentation. Each step feeds the next, with the system managing the entire flow.

06 AUTOMATION


REMOVING THE REPETITIVE FROM WORK



Traditional automation follows fixed rules: if X happens, do Y. This works well for predictable, structured processes but breaks down when inputs are variable or judgment is required. AI-powered automation can handle variability it can read an email, understand its intent, categorise it, draft a response, and route it to the right person.

CATEGORIES OF AUTOMATABLE WORK

Category	Description
Data processing	Extracts, transforms, and routes information from documents, emails and databases.
Communication	Drafts and sends standard responses, follow-ups and notifications.
Scheduling	Manages calendars, books meetings and sets reminders.
Reporting	Aggregates data from multiple sources and generates regular status reports.
Content generation	Produces first drafts of product descriptions, social posts and summaries.
Quality review	Checks documents, code or data against defined standards and flags issues.



A PRACTICAL WORKFLOW EXAMPLE

A freelancer receiving client inquiries by email could build an AI automation workflow that: monitors the inbox for new messages identifies patterns of a new inquiry extracts key details (name, project type, budget, timeline) generates a personalised response from a template logs the inquiry in a spreadsheet and flags it for follow up.

STARTING POINT

Find one repetitive task in your own work something you do the same way every time and ask: what information does this task require, and what does it produce? That structure is your automation blueprint.

07 AI VIDEO MAKING

CREATING VISUAL CONTENT WITH AI



Video has long required a significant production apparatus: cameras, lighting, editing software, voiceover artists, and motion designers. AI is systematically reducing the barriers to each of these steps, making video creation accessible to individuals and small teams who previously could not produce it at scale.

THE SCRIPT-TO-VIDEO PIPELINE


The most common AI video workflow starts with a script and ends with a finished video. AI can now assist or fully handle every step: Script (written with or by an LLM) Voiceover (text-to-speech AI, often indistinguishable from human narration) Visuals (AI-generated imagery or avatar presenters) Timing (audio and visual elements synchronised automatically) Editing (cuts and pacing applied through AI-assisted tools) Captions (automatically transcribed and styled).

USE CASES BY AUDIENCE


Audience	Use Case
 Creators	Create explainer videos, social content and course materials without hiring a production team.
 Educators	Create lecture recordings, concept walkthroughs and visual summaries of written content.
 Businesses	Create product demos, onboarding videos, internal communications and marketing content.
 Developers	Create tool documentation, demo presentations and screencasts with AI narration.

AI video tools are good at execution. They are not yet good at creative direction, brand voice, or knowing what a specific audience actually responds to. A human still needs to define the goal, review the output, and make judgment calls about what works.


AI VIDEO WORKFLOW



USE CASES



YouTube Explainers



Training Videos



Marketing Content



Social Media Videos

KEY INSIGHT

Do not start with the most complex video you want to make. Start with the simplest: a 60-second explainer with a written script, AI voiceover and simple visuals. Build your intuition for what the tools can do before expanding the scope.

08 VIBE CODING

BUILDING SOFTWARE WITH AI COLLABORATION



Vibe coding is the practice of building software through conversation with AI, where the developer describes what they want in plain language and the AI produces working code iterating through natural dialogue rather than formal specification.

A traditional development workflow involves writing code, running it, reading errors, debugging, and repeating. Vibe coding compresses this loop by using AI as an active partner throughout. You describe what you are building, the AI writes the first version, you describe what needs to change, the AI revises and so on until the thing works.

WHAT VIBE CODING ENABLES

- ◆ **Rapid prototyping:** Ideas can become working demos in hours rather than days.
- ◆ **Cross-domain building:** Marketers, designers, and writers can build functional tools without committing to learning to code.
- ◆ **Faster iteration:** Changing direction does not require rewriting everything from scratch.

WHAT IT DOES NOT REPLACE

Vibe coding is well-suited to small tools, internal utilities, prototypes, scripts, and simple web applications. It is less suited to complex systems requiring careful architecture, security-critical code needing formal review or large collaborative codebases where structure and consistency matter deeply.

AI GENERATES

```
1 <form class="login">
2   <input type="email"
3     placeholder="Email">
4   <input type="password">
5     placeholder="Password">
6   <a href="#">
7     Reset Password</a>
8   <button>Login</button>
8 </form>
```

KEY INSIGHT

Vibe coding works best when you treat the AI as a capable but junior developer who needs clear direction. The more precisely you describe desired behaviour inputs, outputs, edge cases, constraints the better the AI performs. Imprecise direction produces imprecise code.

09 PYTHON


THE LANGUAGE THAT POWERS AI



Python has become the defactor language of the AI field not because it is the fastest or most technically elegant, but because it combines readability, flexibility and an extraordinary ecosystem of libraries. Nearly every major AI framework, automation tool and data processing system either runs on Python or has strong Python support.

WHERE PYTHON SHOWS UP

Domain	Libraries & Tools
 Machine Learning	scikit-learn, PyTorch — standard tools for training and evaluating models.
 LLM Integration	APIs for GPT, Claude, and Gemini are all accessible via Python with a few lines of code.
 Automation	Scripts for file handling, email, web scraping, and system tasks.
 Data Processing	Pandas and NumPy for cleaning, analysing, and transforming datasets.
 Web Development	FastAPI and Django power backend systems and APIs.
 AI Agents	LangChain, CrewAI, and AutoGen — leading agent frameworks — are built in Python.



PYTHON IN AN AI WORKFLOW

Python reads more like English than most programming languages. A line like ‘for item in list: print(item)’ is almost self-explanatory. A practical AI workflow might: fetch data from an API clean and structure it with pandas send it to an LLM via the provider’s Python API parse the response and write output to a file schedule the script to run automatically at set intervals.



AI & ML Models



Data Analysis



Automation Scripts



APIs & Integrations

KEY INSIGHT

The best way to learn Python is by building something real. Instead of studying concepts without context, choose a small problem you care about solving and learn the required Python skills along the way.

CORRELATION & VISUALISATION



Correlation measures how strongly two variables relate to each other, ranging from -1 (strong negative) to $+1$ (strong positive). Study hours and exam scores might correlate at $+0.85$; absenteeism and scores at -0.72 . A heatmap is the most efficient way to visualise correlations across all columns at once.






Chart Type	Purpose
 Histogram	Shows the distribution of a single numerical variable.
 Box Plot	Displays median, quartiles, and outliers in one compact visual.
 Scatter Plot	Reveals the relationship between two numerical variables.
 Bar Chart	Compares counts or values across categories (e.g., sales by region).
 Heatmap	Visualises correlations or missing value patterns across the full dataset.

Data cleaning is the process of identifying and correcting errors. It can account for up to 60–80% of a data scientist’s total project time. Key tasks include: handling missing values (by deletion, mean/median imputation, or advanced KNN methods); detecting and treating outliers (using Z-score or IQR methods not all outliers are errors); scaling numerical features (Min-Max Scaling or Standard Scaling) so that large-range columns do not dominate smaller ones; and encoding categorical variables (Label Encoding for ordered categories, One-Hot Encoding for unordered) into numbers models can process.

FEATURE ENGINEERING & PIPELINE DESIGN



Feature engineering creates new input variables from existing data such as extracting weekday from a date column, or creating a price-per-sqft field that improved one house-price model's accuracy from 76% to 88%. Feature selection then removes irrelevant or redundant variables to reduce model complexity and prevent overfitting.

A data pipeline chains all preparation steps cleaning, scaling, encoding and feature engineering — into one repeatable automated workflow. In Python, scikit-learn's Pipeline class makes building and applying pipelines straightforward and saving the pipeline with joblib enables consistent deployment.

KEY INSIGHT

Data preparation is not just a step — it is the foundation of every successful machine learning project. Mastering EDA, Data Cleaning, Feature Engineering and Pipeline Design gives you the toolkit to work with real world messy data confidently.


11 MACHINE LEARNING

CLASSIC ALGORITHMS, EVALUATION AND DEPLOYMENT



Machine Learning (ML) is a branch of AI that enables computers to learn from data and make decisions without explicit programming. Instead of following fixed rules, it identifies patterns from examples. ML powers spam filters, voice assistants, fraud detection, and product recommendations. Its main types are Supervised, Unsupervised and Reinforcement Learning.

ML LIFECYCLE



TYPES OF MACHINE LEARNING

Type	Description
 Supervised Learning	Learns from labelled input-output pairs. Used for classification and regression.
 Unsupervised Learning	Finds hidden patterns in unlabelled data. Used for clustering and dimensionality reduction.
 Semi-Supervised	Uses a mix of labelled and unlabelled data to improve learning efficiency.
 Reinforcement Learning	An agent learns by interacting with an environment and receiving rewards or penalties.
 Self-Supervised	The model creates its own labels from raw data (e.g., GPT, BERT).



THE ML WORKFLOW



Every ML project follows a clear sequence: (1) Define the Problem (2) Collect & Clean Data (3) Feature Engineering (4) Choose & Train the Model (5) Evaluate & Improve (6) Deploy. Classic ML algorithms Linear Regression, Logistic Regression, Decision Trees, k-NN, and SVM are interpretable, data-efficient, and fast to train, making them the ideal foundation before advancing to deep learning.

KEY SUPERVISED LEARNING ALGORITHMS

Algorithm	Key Characteristics
Linear Regression	Predicts continuous values by fitting the best straight line. Fast and highly interpretable.
Logistic Regression	Classifies outcomes using a sigmoid function (e.g., Spam/Not Spam). Great baseline.
Decision Tree	Makes yes/no splits to reach a decision. Intuitive but prone to overfitting.
k-NN	Classifies by majority vote from the k nearest data points. No training phase needed.
SVM	Finds the optimal boundary between classes; powerful for high-dimensional data.

UNSUPERVISED LEARNING

Unsupervised learning finds hidden patterns in unlabelled data. K-Means groups data into k clusters based on distance from cluster centres (the optimal k is found using the Elbow Method). DBSCAN identifies clusters of any shape and automatically labels outliers as noise ideal for geographic data. PCA (Principal Component Analysis) reduces the number of features while preserving the most important information, speeding up training and enabling visualisation of high dimensional datasets.

ENSEMBLE METHODS

Ensemble methods combine multiple models for more accurate and reliable predictions. Bagging (e.g., Random Forest) trains models on random data subsets in parallel and combines them by majority vote or averaging. Boosting (e.g., XGBoost, LightGBM) trains models sequentially, each correcting the errors of the previous one. XGBoost has won hundreds of ML competitions on structured data. Stacking uses a meta-learner to combine the predictions of multiple diverse base models.

MODEL EVALUATION



Data is split into training (80%) and test (20%) sets — never evaluate on the same data used to train. K-Fold Cross-Validation averages results across k different splits for more reliable accuracy estimates. Two critical failure modes: underfitting (model too simple, high bias) and overfitting (model memorises training data, high variance). Overfitting is reduced with regularisation, more data, simpler models, or cross-validation.

Metric	Used For	What It Measures
Accuracy	Classification	Percentage of correct predictions — misleading for imbalanced classes.
Precision / Recall	Classification	Precision avoids false alarms; Recall avoids missed cases. F1 balances both.
ROC-AUC	Classification	The model's ability to distinguish between classes.
MAE / RMSE	Regression	Average prediction error — RMSE penalises large errors more heavily.
R ² (R-squared)	Regression	Proportion of variance explained by the model (1.0 = perfect fit).

MODEL DEPLOYMENT

Deployment makes a trained model available via an API. A REST API (built with Flask or FastAPI) allows clients to send data and receive predictions as JSON. Batch prediction runs the model on large datasets overnight. Model monitoring tracks performance in production, detecting data drift and model degradation over time. Tools like MLflow, BentoML, and Hugging Face make packaging and deploying models easier.

REAL WORLD EXAMPLES



Recommendation Systems



Fraud Detection



Demand Forecasting



Image Recognition

KEY INSIGHT

Mastering Classic ML from supervised and unsupervised learning through ensemble methods and model evaluation gives any aspiring data scientist the toolkit to build reliable, explainable, and deployable machine learning systems.

THE FUTURE OF AI



Predicting the trajectory of AI with precision is difficult the field has repeatedly surprised observers in both directions. Four trends deserve attention.



Multimodal Systems: AI is increasingly capable across text, images, audio, video, and code simultaneously blurring the lines between specialised tools and expanding single-workflow possibilities.



Smaller, More Efficient Models: The race toward ever-larger models is being complemented by significant investment in smaller, faster models that run locally or at lower cost making capable AI more accessible.



Agent Networks: As individual AI agents mature, coordinated networks of specialised agents handling different parts of complex workflows represent the next meaningful evolution of agentic AI.




AI in Physical Systems: Robotics and autonomous systems are beginning to incorporate modern AI reasoning. The gap between digital AI and physical-world AI is narrowing, with implications for logistics, manufacturing, and services.

ETHICS AND RESPONSIBLE USE

As AI becomes more embedded in consequential decisions, several considerations deserve ongoing attention:

- verify factual claims from AI
- hallucination is real; be aware of bias
- AI reflects the data it was trained on;
- understand transparency knowing roughly how an AI reaches its outputs matters for appropriate trust; and review data privacy
- using AI tools involves sharing data with third-party systems.







Verify Outputs



Protect Privacy



Use Judgement



Stay Informed

KEY INSIGHT

Mastering Classic ML from supervised and unsupervised learning through ensemble methods and model evaluation gives any aspiring data scientist the toolkit to build reliable, explainable, and deployable machine learning systems.

HANDS-ON ACTIVITIES



Understanding AI conceptually is the start, not the end. These activities are designed to move you from passive reading to active experimentation. Most require only a freely available AI tool and a curious mindset.

ACTIVITY 1 — THE PROMPT COMPARISON LAB

Write the simplest possible prompt on any topic you know well. Then write a second prompt that includes your background, what you specifically want, the format, and any constraints. Compare the two outputs side by side. Which information made the most difference?


Prompt Playground

Try weak vs strong prompts.



Agent Designer

Design your own AI agent workflow.



ACTIVITY 2 — MAP AN AGENT WORKFLOW

Pick a multi-step task from your own work. On paper, break it into the individual steps an AI agent would take. For each step, identify what information the agent needs, what action it takes, and what it produces. Where is the workflow most fragile?

ACTIVITY 3 — AUTOMATION BLUEPRINT

Identify one task you do at least weekly that follows a consistent pattern. Write out every sub-step, sketch a flow (trigger steps output), and label which parts are AI, automated logic, or manual. What would you need to change to make this workflow possible?

Automation Mapper

Map a task you want to automate.



Vibe Coding Challenge

Build a small app with AI help.



ACTIVITY 4 — VIBE CODE SOMETHING SMALL

Identify a small, specific problem: a calculator for something you do manually, a word counter, or a list formatter. Describe the problem to an AI coding assistant in plain language. Test the output; if something does not work, describe what went wrong and ask for a fix.

ACTIVITY 5 — EDA MINI-PROJECT

Find a freely available dataset (Kaggle, government open data). Generate basic statistics — min, max, mean, missing values. Create a histogram and a scatter plot. Write a one-paragraph summary of what you discovered. What cleaning would this data need?

Data Explorer

Analyze a dataset using visuals and insights.



ML Thinker

Predict outcomes from simple data patterns.



ACTIVITY 6 — ML ALGORITHM SELECTION CHALLENGE

For each scenario, choose an algorithm and explain why: (a) Predict house prices from size and location. (b) Group customers by purchase history. (c) Classify emails as spam or not spam. (d) Detect unusual financial transactions. What factors drove your choices?

FINAL SUMMARY


Chapter	Core Takeaway
AI Foundations	AI systems learn patterns from data and use those patterns to make useful predictions. Sophisticated but not magic.
Prompt Engineering	How you communicate with AI shapes what you get. Precision, context and iteration are the core skills.
Large Language Models	LLMs generate text by predicting what comes next. Powerful within limits and those limits matter as much as the capabilities.
AI Agents	Agents use LLM reasoning to take action and manage workflows not just respond to individual prompts.
Agentic AI	Agentic systems operate autonomously across complex multi-step goals beginning to change how knowledge work is organized.
Automation	AI expands automation beyond rule based tasks into anything involving language judgment or pattern recognition.
AI Video Making	AI has made video production accessible to individuals and small teams by assisting every step of the pipeline.
Vibe Coding	Building software through AI conversation is a defining shift in how software gets made and who can make it.
Python	Python ties AI automation and data work together Literacy not mastery opens the right doors.
EDA	EDA is how you understand clean and prepare data before modeling It is the foundation of every reliable ML project.
Machine Learning	Classic ML algorithms provide the foundation for making predictions finding patterns and deploying intelligent systems.

THE BIGGER PICTURE

These eleven topics are not isolated skills. They form a stack: Foundations support Prompting, Prompting makes LLMs useful, LLMs power Agents, Agents enable Agentic Systems, Agentic Systems drive Automation, Automation creates capacity for creative work like Video, Vibe Coding lowers the barrier to building, Python provides the technical substrate, EDA transforms raw data into trustworthy inputs and ML turns those inputs into actionable intelligence.

Understanding that stack even at the introductory level this guide provides puts you in a meaningfully different position than someone who knows these topics only by name.





WHERE TO GO FROM HERE

The natural next step is depth. Pick the topic that felt most relevant to your actual situation and spend real time with it not just reading about it, but doing something with it. The difference between someone who has read about AI and someone who has worked with it is significant, and that gap closes faster than most people expect.



9676125666 , 9676327666



LEARN_AIATOZ



WWW.AIATOZ.ORG



LEARN_AIATOZ



WWW.LINKEDIN.COM/IN/AIATOZ



LEARN_AIATOZ



AIATOZOFFICIAL@GMAIL.COM